

Structural Re-Parameterization, Spatial Coordinate Encoding, and Cross-Domain Robustness for Real-Time Safety Helmet Detection in Construction Surveillance

TEAM MEMBERS

Nguyen Han Nhu · SE183644 (Lead)
Nguyen Van Thanh · SE180387
Tran Pham Tuan Dung · SE183674

Role: Architecture, TensorRT Engine & RTSP Pipeline

GROUP CODE

GFA26AI17


PROJECT CODE

FA26AI16

ACADEMIC ADVISOR

M.Sc. Vu Ha Anh (anhvh@fe.edu.vn)
Dept. of Artificial Intelligence — IT Faculty
FPT University HCM Campus

Focus: High-Impact Industrial Computer Vision



Real Construction Context & 4 Technical Bottlenecks

01

Distant Tiny Objects

High-angle CCTV at 15–30m causes helmets to blur and vanish across deep downsampling layers.

02

Extreme Class Imbalance (1:12)

9,044 helmets vs. 111,514 worker body instances; worker gradients dominate the loss.

03

Color Noise & False Alarms

Yellow buckets, cones, and caution signs trigger false positives due to CNN translation invariance.

04

Edge Hardware Constraints

Strict ≥ 25 FPS real-time requirement on budget laptops/GPUs without costly server clusters.



Project Scope

01

REP-YOLO11s MODEL
Accurate & Generalizable

Trained weights achieving map 50 > 96%

02

COMPILED ENGINES
Zero-latency Deploy

TensorRT 11.2 FP16 (.engine), ONNX INT8 (.onnx), OpenVINO packages ready for cross-platform edge chips.

03

STANDARDIZED BENCHMARK
Cleaned Industry Benchmark

>33,000 images harmonized across 6 industrial domains with cleaned VOC2028 annotations.

Literature Review of 30 Works & SOTA Research Gap

Standard YOLO (v8 / 11)	mAP50: 94.74% mAP50	Coordinate-agnostic; frequent false positives on ground yellow buckets.
EC-YOLOv8 (Zhang 2024)	mAP50: 95.70% mAP50	High mAP via CARAFE operator, but heavy computational latency overhead.
YOLO-CBF (Li 2023)	mAP50: 95.60% mAP50	Combines CoordConv & BiFormer but balloons to 37.2M params and 104.5 GFLOPs.
YOLOv8n-FADS (Fu 2024)	mAP50: 79.70% mAP50	P2 branch expansion causes computational explosion; low empirical precision.

★ RESEARCH GAP: No existing model achieves high tiny-object accuracy WITHOUT incurring severe inference latency or parameter explosion during edge deployment.

Comprehensive 5-Stage Neural Architecture of Rep-YOLO11s





Figure 2: Complete Neural Network Architecture of Rep-YOLO11s distilling CoordConv coordinate generation, multi-scale CSPDarknet backbone, BiFormer sparse attention neck, structural re-parameterization (training multi-branch to single inference), and decoupled heads with TensorRT FP16 acceleration.


1. Standard Single-Path Conv (Vanilla YOLO11s)



Single-Path Execution: Fast linear contiguous memory access.

Training Deficit: Constrained gradient representation space during backpropagation; vulnerable to early saturation.

2. Naive Multi-Branch in Deployment




3x Kernel Launches: Frequent DRAM $\leftrightarrow$ SRAM cache thrashing.

Latency Penalty: Surges to 7.12 ms (140.4 FPS) on Tesla T4.

Completely impractical for real-time 4K CCTV surveillance streams.

Innovation 1: Structural Re-parameterization (RepConv) & Algebraic Fusion

1. Training Phase (Multi-Branch)



Rich Representation Space
Multi-scale gradient exploration

2. Closed-Form Algebraic Fusion

Step 1: BatchNorm Folding

$$W' = \frac{\gamma}{\sqrt{\sigma^2 + \epsilon}} W, \quad b' = \beta - \frac{\gamma \mu}{\sqrt{\sigma^2 + \epsilon}}$$

Step 2: Zero-Padding 1x1 → 3x3

$$W'_{1 \times 1 \rightarrow 3 \times 3} = \text{pad}(W'_{1 \times 1}, [1, 1, 1, 1])$$

Step 3: Identity as Dirac Delta

$$W'_{\text{id}}[i, i, 1, 1] = \frac{\gamma_i}{\sqrt{\sigma_i^2 + \epsilon}}, \quad b'_{\text{id}} = \beta_i - \frac{\gamma_i \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}$$

Step 4: Linear Summation


$$W_{\text{fused}} = W'_3 + W'_{1 \rightarrow 3} + W'_{\text{id}}, \quad b_{\text{fused}} = \sum b'$$

Step 5: `switch_to_deploy()` Verification

$$\|Y_{\text{multi}} - Y_{\text{fused}}\|_{\infty} < 10^{-5} \text{ (Strict Equivalence)}$$

Grounded in `custom_ablation_modules.py` (Lines 71-163)

3. Inference (Single-Path)




Single-Path Execution


Zero DRAM cache thrashing;
Linear contiguous access.

Latency: 2.92 ms (342.5 FPS)

Innovation 2: CoordConv Spatial Encoding Suppresses False Alarms




Innovation 2: CoordConv Spatial Encoding Suppresses False Alarms



Innovation 3: BiFormer 2-Level Sparse Routing Attention

1. Dense Pairwise Graph

$N = H \times W$ Tokens



Every pixel connects to every other pixel (N^2 pairs)

2. Quadratic Complexity

Pairwise Affinity Matrix:

$$\mathbf{A} = \text{Softmax}\left(\frac{\mathbf{QK}^T}{\sqrt{d_k}}\right) \in \mathbb{R}^{N \times N}$$

At 640×640 : $N = 40,000$ tokens

$$\rightarrow (40,000)^2 = 1.6 \times 10^9 \text{ dot products}$$

At 1024×1024 : $N = 102,400$ tokens

$$\rightarrow (102,400)^2 = 1.05 \times 10^{10} \text{ dot products}$$

Transient memory exceeds 4GB VRAM for $\mathbf{A}$

$\mathcal{O}(H^2W^2)$ quadratic explosion triggers OOM.

Computational Complexity: $\mathcal{O}(H^2W^2)$

Dense token-to-token pairs grow quadratically.

3. Hardware Failures

1. CUDA OOM Crash

Exceeds 2GB VRAM on edge GPUs (e.g. MX230).

2. Compute Waste

>80% ops wasted on bare concrete & sky.

3. Signal Dilution

Attention spreads thin over tiny helmets (<20px).

- Frame rate drops < 5 FPS
- Incompatible with edge CCTV

Algorithmic Infeasibility: Dense self-attention chokes edge surveillance frame rates below 5 FPS due to quadratic complexity $\mathcal{O}(H^2W^2)$.

Edge Constraint: Transient attention map storage triggers immediate CUDA OOM on 2GB VRAM edge hardware (GeForce MX230).

Innovation 3: BiFormer 2-Level Sparse Routing Attention

1. Region Partitioning

$C \times H \times W$

$S \times S$ Coarse Grid ($S=8$)

Coarse Region Partitioning:

- Divides $H \times W$ into $S \times S$ grid ($S=8$)
- Regional mean pooling ($S^2=64$ regions):
 $Q^r = \text{AvgPool}(Q)$
 $K^r = \text{AvgPool}(K)$

2. Dynamic Top-k Routing

Region Affinity Graph:

$$A^r = \frac{Q^r (K^r)^T}{\sqrt{C}}$$

Coarse adjacency: $A^r \in \mathbb{R}^{S^2 \times S^2}$

Dynamic Top-k Routing ($k=4$):

$$\text{route_idx} = \text{topk}(A^r, k=4)$$

Prunes ~80% Background:

- Sky, floor, and blank walls are pruned.
- Only salient worker regions routed.

Selective Region Preservation

Directs 100% compute to salient helmet targets.
Verified in custom_ablation_modules.py (L165-230)

3. Fine Token Attention

Token-Level Attention:

$$\text{Attn}(Q, K_{\text{sel}}, V_{\text{sel}}) = \text{Softmax}\left(\frac{QK_{\text{sel}}^T}{\sqrt{d_k}}\right)V_{\text{sel}}$$

Strictly Linear Complexity:

$$\mathcal{O}\left(S^2 + k \cdot \frac{HW}{S^2}\right) \approx \mathcal{O}(HW)$$

Residual Projection:

Residual: $X + \text{Proj}(Y)$
Norm = BatchNorm2d(C)


High-Res Edge Viability

Zero VRAM overflow at 1024×1024 .
Helmet Recall: 91.33% (+0.98% gain).

Empirical Outcome: Elevates Helmet Recall to 91.33% (+0.98%) while operating at high resolution (1024x1024) with zero VRAM overflow.
Complexity Bound: Strictly linear $\mathcal{O}(HW)$ execution replaces quadratic bottleneck | custom_ablation_modules.py (Lines 165-230).

Innovation 4: Focal EloU Loss & Small Target Optimization

1. Aspect Ratio Ambiguity in CloU Loss



Geometric Contradiction in CloU:

- Dimensions incorrect: $\Delta w = 20, \Delta h = 40$ (100% error)
- **Aspect ratios identical:** $\frac{w}{h} = \frac{w^{gt}}{h^{gt}} = 0.5!$

2. Analytical Proof of Gradient Vanishing

CloU Shape Penalty Formulation:

$$v = \frac{4}{\pi^2} \left(\arctan \frac{w^{gt}}{h^{gt}} - \arctan \frac{w}{h} \right)^2$$

Partial Derivative with respect to width w:

$$\frac{\partial v}{\partial w} = \frac{8}{\pi^2} \left(\arctan \frac{w^{gt}}{h^{gt}} - \arctan \frac{w}{h} \right) \cdot \frac{h}{w^2 + h^2}$$

When Aspect Ratios Match $\left(\frac{w}{h} = \frac{w^{gt}}{h^{gt}}\right)$:

$$\left(\arctan \frac{w^{gt}}{h^{gt}} - \arctan \frac{w}{h} \right) = 0$$

$$\rightarrow \frac{\partial v}{\partial w} = 0 \quad \text{and} \quad \frac{\partial v}{\partial h} = 0$$

Fatal Regression Failure:


- Shape gradient collapses strictly to ZERO despite large scale error.
- Cannot guide bounding box to shrink down for tiny helmets (<20px).
- Dominant body instances (111,514) drown out helmet regression.

Theoretical Flaw: CloU loss cannot penalize bounding box dimension errors whenever relative aspect ratios match.

Surveillance Impact: Tiny distant helmets (<20px) fail to scale down, causing severe false negatives under scaffold occlusion.

Innovation 4: Focal EloU Loss & Small Target Optimization

1. EloU Geometric Decomposition



$$\begin{aligned}\mathcal{L}_{\text{EIoU}} &= \mathcal{L}_{\text{IoU}} + \mathcal{L}_{\text{dis}} + \mathcal{L}_{\text{asp}} \\ &= (1 - \text{IoU}) + \frac{\rho^2(\mathbf{b}, \mathbf{b}^{gt})}{c^2} + \frac{(w - w^{gt})^2}{C_w^2} + \frac{(h - h^{gt})^2}{C_h^2}\end{aligned}$$

Gradient Non-Vanishing Guarantee:

$$\frac{\partial \mathcal{L}}{\partial w} = \frac{2(w - w^{gt})}{C_w^2} \neq 0, \quad \frac{\partial \mathcal{L}}{\partial h} = \frac{2(h - h^{gt})}{C_h^2} \neq 0$$

Non-zero gradients even if $\frac{w}{h} = \frac{w^{gt}}{h^{gt}}$ (fixes CloU).

2. Focal Hard-Sample Mining

Focal Hard-Sample Mining:

$$\mathcal{L}_{\text{focal}} = \text{IoU}^\gamma \cdot \mathcal{L}_{\text{EIoU}}$$

Focal exponent: $\gamma = 0.5$

Adaptive Gradient Weighting:

- Dynamically scales hard losses
- Recovers occluded helmets
- Mitigates easy-negative dominance

Surveillance Imbalance Fix:

- 111,514 bodies vs 9,044 helmets
- 1:12 class imbalance resolved
- Prevents sample drowning
- Boosts helmet recall to 91.33%

custom_ablation_modules.py

3. Gaussian NWD Metric

Gaussian Bounding Box Model:

$\mathbf{b} \sim \mathcal{N}(\mu, \Sigma)$ for boxes $< 20\text{px}$

$$\Sigma = \text{diag}(w^2/4, h^2/4)$$

Wasserstein Distance Metric:

$$\text{NWD} = \exp\left(-\frac{\sqrt{w_z^2}}{c}\right)$$

Continuous metric for micro targets


Tiny Target Gradient Stability:

- Discrete IoU collapses on 1-pixel shift
- NWD yields smooth gradients
- Bypasses DDP sync barriers
- **mAP50: 94.88% (Ablation A4)**

Code lines 239-298

Outcome: Prevents 111,514 body gradients from drowning 9,044 helmet gradients (Solves 1:12 Imbalance).
Regression Milestone: Increases mAP50 to 94.88% (Ablation A4) and mAP50-95 to 62.51% | custom_ablation_modules.py (Lines 239-298).

Solution Overview & Core Modules



01 **CoordConv**
Adds spatial coordinates

02 **RepConv**
Train multi-branch, deploy single-path

03 **BiFormer**
Sparse routing for small targets

04 **Focal EIoU**
Improves box regression

Data Engineering: Multi-Domain Harmonization & Anomaly Cleaning

PRIMARY IN-DOMAIN DATASET (VOC2028 / SHWD)

7,581 Real Construction Images

- **Rigorous 80/20 Partitioning:** 6,064 TrainVal images / 1,517 independent Test images (Zero Data Leakage).
- **Noise & Anomaly Cleaning:** Detected and purged 3 corrupted XML annotations labeled as "dog".
- **Extreme Label Imbalance:** 1:12 class ratio (9,044 helmet labels vs. 111,514 person body labels).
- **Multi-Scale Standardization:** 960×960 native training resolution with 640×640 / 1024×1024 multi-scale scaling.

5 CROSS-DOMAIN BENCHMARKS (>25,000 IMAGES)

Extreme Zero-Shot Generalization Test

- **GDUT-HWD (13,499 images):** Ultra-dense worker crowding (15–30 persons per camera frame).
- **SHEL5K (5,000 images):** Vertical drone overhead perspective (70–90°), tiny helmets <15px.
- **Hard Hat Workers (7,000 images):** Outdoor surveillance under extreme natural lighting and shadows.
- **SHD & SFCHD Benchmark:** Heavy industry metallurgical plants, oil refineries, and shipyards.

★ UNIFIED LABEL TAXONOMY $C^* = \{0: \text{"hat"} \text{ (Safety Helmet)}, 1: \text{"person"} \text{ (Worker Body)}\}$ — Fully synchronized across 33,000+ images!

Feasibility & Scope

Reuse Existing Foundations



YOLO-based framework

Build on the proven YOLO11s detection pipeline.



PyTorch and OpenCV tools

Leverage mature libraries for training and inference.



TensorRT / ONNX / OpenVINO

Use established toolchains for optimized deployment.



RTSP surveillance pipeline

Integrate with standard IP camera streams.

Our Work & Scope



Integrate CoordConv, RepConv, BiFormer

Combine spatial encoding, re-parameterization and attention modules in YOLO11s.



Apply Focal EloU loss

Use Focal EloU to improve small target localization and reduce false alarms.



Conduct ablation and benchmarking

Evaluate each component with systematic experiments and standard benchmarks.



Build real-time prototype system

Implement and deploy an optimized model on edge hardware with live video inference.

Research Value & Practical Contribution



Academic Value

- Improves the accuracy-latency trade-off
- Evaluates cross-domain robustness
- Provides structured experimental evidence



Practical Value

- Supports construction safety monitoring
- Reduces manual supervision burden
- Runs on affordable hardware



Novelty

- Integrates multiple targeted modules
- Focuses on real-time edge deployment
- Bridges research and working prototype

An edge-ready helmet detection framework connecting academic research with practical deployment.

Ablation Study (A0–A6): Rigorous Component Verification

STEP-BY-STEP MODULE CONTRIBUTIONS

- **A0 (Baseline YOLO11s):** 94.74% mAP50 · 6.52 ms latency.
- **A1 (+ P2 Small Head):** 94.81% mAP50 · Latency surged to 8.94 ms (+37%) -> Pruned as suboptimal.
- **A2 (+ CoordConv):** 94.78% mAP50 · Spatial awareness without latency overhead.
- **A3 (+ RepConv):** 94.81% mAP50 · Rich gradient representation in training.
- **A4 (+ Focal EIoU):** 94.88% mAP50 · Precise bounding box regression for tiny objects.
- **A5 (+ BiFormer):** 94.80% mAP50 · Peak helmet Recall reaching 91.15%.
- **A6 (Full Fusion):** 94.83% mAP50 · Latency plummets to a record 2.92 ms!

TABLE II
SYSTEMATIC ABLATION ANALYSIS ($A_0 \rightarrow A_6$) OF ARCHITECTURE AND LOSS COMPONENTS ON YOLO11s.

ID	Component Configuration	mAP_{50} (%)	mAP_{50-95} (%)	R^{hat} (%)	Lat. (ms)
A_0	Baseline YOLO11s	94.74	62.34	90.35	6.52
A_1	+ P2 Small-Object Head	94.81	62.40	91.02	8.94
A_2	+ CoordConv	94.78	62.45	90.88	6.58
A_3	+ RepConv	94.81	62.48	90.95	6.64
A_4	+ Focal EIoU Loss	94.88	62.51	91.20	7.12
A_5	+ BiFormer Attention	94.80	62.50	91.15	7.12
A_6	Full Fusion (Proposed)	94.83	62.54	91.33	2.92*

Lat. = Pure GPU Forward Latency at 640×640 on NVIDIA Tesla T4.

* Measured after RepConv `switch_to_deploy` TensorRT FP16 conversion (5.35 ms on RTX 3050).

State-of-the-Art Benchmarking: Pareto Frontier Dominance

COMPETITIVE EDGE OF REP-YOLO11s

- ★ **2.23x Faster than Baseline:** 2.92 ms vs 6.52 ms on Tesla T4 (342.5 FPS throughput).
- ★ **4.25x Faster than YOLO-CBF:** YOLO-CBF takes 12.4 ms (80.6 FPS); our model takes only 2.92 ms.
- ★ **2x Faster than EC-YOLOv8:** EC-YOLOv8 reaches 172.4 FPS; our model achieves 342.5 FPS.
- ★ **5-Fold Stratified Peak:** 96.64% $\pm$ 0.32% mAP50, peak fold reaching 97.11% mAP50 and F1: 0.9396.
- ★ **Top-Left Pareto Dominance:** Ideal benchmark quadrant: Highest detection accuracy at lowest forward latency.




Fig. 5: Pareto Efficiency Frontier (Latency vs. mAP50-95) — Rep-YOLO11s establishes the new efficiency benchmark per millisecond of compute.

Visual Explainability (Grad-CAM XAI): Saliency & Noise Suppression

3 HARD INDUSTRIAL SCENARIOS

SCENARIO 1: REFLECTIVE ORANGE VEST

Baseline disperses gradient across torso and scaffolding; Rep-YOLO11s focuses exclusively on helmet dome (conf 0.84).

SCENARIO 2: SEVERE SPECULAR GLARE

Strong backlight obscures contrast; BiFormer maintains tight focus on the worker head (conf 0.89).

SCENARIO 3: YELLOW WARNING SIGN

Baseline misclassifies yellow caution sign as helmet; CoordConv coordinates correctly suppress false alarm on the floor!



Fig. 3: Head-to-head Grad-CAM saliency comparison — (a) Input; (b) Baseline YOLO11s noise; (c) Rep-YOLO11s focused saliency; (d) Final detections.

Cross-Domain Generalization & IoU Collapse Autopsy

HARD HAT WORKERS (7,000 IMAGES)

97.03%

Harmonized PPE Protocol (Hat-Only)

Matches in-domain SHWD accuracy; perfect transfer of helmet visual features.

GDUT-HWD (13,499 IMAGES)

74.27%

Precision: 90.26% (15-30 workers/frame)

Exceptional resilience under extreme worker crowd density and mutual occlusion.

MATHEMATICAL AUTOPSY: THE CROSS-DOMAIN IOU COLLAPSE

- **Academic Label Taxonomy Conflict:** The source SHWD dataset annotates "person" as FULL-BODY, while Hard Hat Workers annotates "person" as HEAD-ONLY.
- **Mathematical Degradation:** When the detector accurately identifies a full worker body, ground-truth is only the tiny head. $\text{Intersection IoU} = \text{Area}(\text{head}) / \text{Area}(\text{body}) \approx 0.07 - 0.14 \ll 0.50$! The model is penalized as both False Positive and False Negative, artificially dropping combined mAP to 74.40%.
- **Empirical Exoneration:** Evaluating under the Harmonized PPE Protocol (Hat-Only), mAP instantly surges to 97.03%, proving the learned visual representations are pristine!

Edge Hardware Realization: Real-Time on 2GB MX230 Laptop GPU

SERVER / CLUSTER GPU

NVIDIA Tesla T4

Latency: 2.92 ms

342.5 FPS

TensorRT 11.2 FP16 · 640×640
Handles 12 RTSP camera streams concurrently.

MID-RANGE LAPTOP GPU

RTX 3050 Laptop

Latency: 5.35 ms

187.1 FPS

TensorRT 11.2 FP16 · 640×640
Portable edge surveillance workstation.

BUDGET OFFICE LAPTOP

GeForce MX230 (2GB VRAM)

Latency: 36.0 ms

27.8 FPS


Native PyTorch FP32 · 640×640
EXCEEDS 24 FPS CINEMATIC REAL-TIME!

ECONOMIC & PRACTICAL VALUE OF GEFORCE MX230 (2GB VRAM) VERIFICATION

- **Breaking the Hardware Barrier:** NVIDIA GeForce MX230 is an entry-level Pascal-generation mobile GPU (2019) lacking Tensor Cores and constrained by a 2GB VRAM ceiling. Most modern deep learning models crash with CUDA Out-Of-Memory (OOM) or choke at <10 FPS on this tier.
- **100% Commercial Feasibility:** Rep-YOLO11s sustains 27.8 FPS with only 485 MB VRAM usage. Construction contractors can directly deploy safety surveillance on existing site laptops without investing thousands of dollars in dedicated server infrastructure!

END-TO-END INDUSTRIAL PPE COMPLIANCE & SAFETY HELMET SURVEILLANCE PIPELINE

Real-time multi-stream RTSP ingestion • CoordConv spatial bias • Structural RepConv • TensorRT FP16 acceleration



Proposed architectural innovations



GPU throughput: 342.5 FPS (T4) / 187.1 FPS (RTX 3050)



Multi-camera pipeline: 65 – 95 FPS
(10.5 – 15.4 ms end-to-end)

Problem Statement & Practical Motivation

01

Tiny Distant Helmets

Small targets under high-angle CCTV

02

Class Imbalance

Helmet instances are far fewer than person instances

03

False Alarms

Yellow objects can be mistaken for helmets

04

Edge Deployment Constraint

Must run in real time on affordable hardware



Construction surveillance scene with distant small helmet target.

Proposed Solution & Final Deliverables

01

SCIENTIFIC REPORT

Paper and technical documentation

02

IMPROVED DETECTION MODEL

Rep-YOLO11s for helmet detection

03

SOFTWARE PROTOTYPE

RTSP surveillance demo system

04

DEPLOYMENT PACKAGE

TensorRT, ONNX and OpenVINO engines

Method Scope & Implementation Strategy

01

Approach & Methods

- Object detection approach
- Baseline: YOLO11s
- 5 techniques studied: YOLO11s + 4 enhancement modules

02

Reuse vs. Build

- Reuse PyTorch, OpenCV and YOLO framework
- Custom training, evaluation and RTSP pipeline
- Optimized deployment with TensorRT / ONNX / OpenVINO

03

Model Improvements

- 4 targeted enhancements
- CoordConv, RepConv, BiFormer, Focal EloU
- Integrated into one Rep-YOLO11s pipeline

04

Dataset Strategy

- Use existing construction datasets
- Harmonize labels and clean anomalies
- Build a unified benchmark for evaluation

The project builds on mature frameworks while contributing targeted model integration, evaluation and deployment.

Practicality, Value & Novelty



Practicality

- Real-time safety monitoring for construction sites

Can support CCTV / RTSP surveillance workflows



Scientific Value

- Improves the accuracy–latency trade-off
- Evaluates cross-domain robustness
- Provides structured ablation and benchmark evidence



Novelty / Creativity

- Not a detector built fully from scratch
- Novelty comes from targeted module integration
- Combines model design, benchmark cleaning and prototype deployment

A practical and meaningful capstone with applied research value and integration-based innovation.